


The For-Profit AI Safety Landscape



When I started thinking about AI Safety (AIS) as a career path in 2023, the shapes of employers that naturally surfaced were the frontier labs, nonprofits (research, fieldbuilding, philanthropy), and policy/governance orgs (think tanks, government). In 2026, I would argue there's another meaningful different-shaped path (which is already a tempting option for many of my friends): for-profit AIS startups are getting funded and are starting to do useful work in the AIS ecosystem.

This is an attempt to make sense of the emerging for-profit AIS landscape. I won't go deep into what makes for-profit incentives useful for AIS, but as a TLDR: I think for-profit companies are well-suited for solving certain narrow problems in AIS, and I think the best strategy for mitigating AI risk is a defense-in-depth approach. (Manifund's "D211 default to non-profit" makes that case well.) This is not an attempt to create an exhaustive list, instead I cover layers where I think for-profit entities can work well in solving concrete problems (and highlight a few that I think are doing good work and/or are finding good fit).

Thanks for reading Yury's Substack!
Subscribe for free to receive new posts and support my work.

Pledge your support

Disclosure: I work for Lionheart Ventures. Lionheart is an investor in Goodfire, Gray Swan, Softmax, Kradle, Lucid Computing, Glacis, and other startups.

→ my arguments were more around incentivizing people to consider this (for-profit as a solution to AIS)

Given that I think startups can be good at solving narrow safety/security problems, I organize them into defense-in-depth layers (s/o to AI Safety, Ethics and Society). I identify seven layers (with some overlap between them), from model development to deployment, then applications, and finally downstream and long-term societal effects.

Training/Research

The earliest safety intervention can be (at least theoretically) done at the root by making the models themselves safer and more interpretable - in the limit, solving alignment and/or making AIS fully transparent and deterministic. Most of the work in training and relevant research is done by the labs, and everything outside carries a high-risk, high-reward profile: new architectures, training methods, interpretability research that can affect the overall direction of model development and/or make progress in alignment and understanding. Though I expect many if not most of these bets will fail, I still think they're worth making - they pull outside funding and new talent into safety research, and some directions (ex. interpretable and narrow superhuman AIS) seem promising conditional on intentional slowdown in capabilities with LLMs.

Highlights:

→ there is both research and startup

• **Goodfire** - An interpretability lab and probably the best-known AIS startup. They build interpretability tools that find unwanted behaviors in models: their amplification method turns backdoors that fire once in 5,000 rollouts into ones that fire once in 50, much better elicitation, and their probes already run in production for enterprise customers, screening personal data out of user messages before a major e-commerce company's models. The long-term vision would be interpretability tools running alongside every deployed model, which gives us a verification layer for AI in production - the ability for anyone to check what a model is actually doing, which barely exists today. Goodfire has hired from Palantir and the major AI labs, and raised \$150M at a \$125B valuation in February 2026, which included top VC firms that don't have a focus on AIS.

• **Softmax** - A research lab betting that alignment can emerge directly from training, with agents learning to care about shared goals the way cells cooperate into organisms. They're building a massively multiplexer benchmark for social intelligence, with open-source environments and training tools for agents to learn that behavior in. Backed by a16z and Lionheart, they again bring in substantial counterfactual funding and hires to support a well-resourced alignment research org.

→ to think about more: what makes a startup explicitly "safety" oriented (does monitoring laborability count, does it have to be related to existential risk, etc)
- there is a lot of money flowing through both regular startup ecosystem and safety right now, so doesn't seem like these startups are doing "safety-washing" and that's interesting so consider

(Pre)-Deployment/Applications

Once a model is trained, the next line of defense is the application layer. First pre-deployment evaluations for capabilities and safety (ideally a close proxy for real-world deployment), and once the model is live, monitoring, steering, and permissioning what it can access. While pre-deployment evals are essential and inform model development, they're not sufficient on their own, as frontier AIS are highly aware of when they're being tested. From what I've seen, this is the broadest, and perhaps most established and competitive market, covering everything from evals to AI application design.

→ perhaps capabilities is only evals and so popular + engineer's skillset might be closer

→ wonder if there will be a downturn in eval companies

Highlights:

• **Andon Labs** - Working toward a "safe autonomous organization": AIS are starting to run businesses and physical systems without humans in the loop, and Andon builds realistic evals of full autonomy ahead of this integration. In their evals, Claude ran real businesses - a vending machine in Anthropic's office (famously tanking its net worth on tungsten cubes), a retail shop and a cafe, and their robotics evals put LLMs in embodied hardware. In July 2026 their arena caught Opus 5 proposing or joining price-fixing cartels in all six runs, while GPT-5.6 Sol refused and asked for its disqualification. By evaluating autonomy in realistic scenarios, Andon creates insights into dangerous propensities such as collusion and deception in more controlled environments, before we notice them in real deployments.

• **Kradle** - Builds multiplayer evaluation worlds where frontier models cooperate, compete, and negotiate against each other over long horizons. Multi-agent systems are where AI applications are heading (frontier labs already use swarms of agents for internal work), making it a new frontier for evals. Kradle's latest research measured how often frontier models lie when lying pays. The team pairs a Firebase co-founder with a former TensorFlow product lead and an alignment researcher from Mila.

• **Apollo Research** - Runs pre-deployment evaluations focused on scheming and deception - whether models strategically hide capabilities or pursue goals they know their developers don't want. Their evals appear in frontier model system cards, and they became a public benefit corporation in January 2026 after operating as a nonprofit since 2023, concluding that demand for evals work couldn't be served at scale on philanthropic funding.

→ how does this become profitable?
→ already Dax' recent push on funding for models → will this be different for safety companies?

Software Security

Calling AI systems software is an oversimplification, but I group security against AI-enabled cyber risk and security of AI systems (ex. adversarial robustness) together because they exist in adjacent market categories, and protect against similar threat models. Companies in this segment cover both reducing an attacker's cyber offensive capabilities and securing the targets. I'm also inclined to group them together as many startups working on red-teaming call themselves AI security startups.

→ me these have a very clear profit/PMF strategy

→ perhaps Outrage belongs here too?

Highlights:

• **Gray Swan** - Does red-teaming and works with every major lab - they run the largest public jailbreaking arena and appear as a security evaluator in frontier model system cards. They're also one of the largest AIS startups currently, having raised a \$40M Series A at a \$200M post-money valuation in May 2026. Their co-founder and Chief Scientist Zico Kolter also chairs the Safety & Security Committee on OpenAI's board.

• **Haize Labs** - Does automated red-teaming - adversarial search that generates attacks at scale instead of waiting for humans to think of them. They red-teamed o1 before launch, and their open-source judges match frontier models at catching failures, at a fraction of the cost - which also makes them usable as reward signals for training the failures out. They raised at a \$100M valuation seven months after founding, straight out of Harvard undergrad.

• **Irregular** - Runs offensive-cyber evals for frontier models. They developed capture-the-flag and full network-attack simulations that appeared in Anthropic's system cards for Claude 3.7 Sonnet and Sonnet 4.5, and in OpenAI's o3 and GPT-5 cards, under their former name Pattern Labs. They raised \$80M from Sequoia and Redpoint at a \$450M valuation.

Securing Hardware

In my opinion, one of the most impactful and best for-profit suited sectors - hardware-level security against malicious actors (both human and AIS) covers access to compute and using it to cause harm. This layer also covers non-proliferation (ex. location verification for chips) and security against AIS enforcment their own weights. Startups here can build the physical enforcement infrastructure that policy will depend on.

→ cautious about if this will be inherently tied to national security/US-China relations

→ specifically reducing dependencies on Taiwan

Highlights:

• **Lucid Computing** - Builds verification infrastructure for AI hardware. The Chip Security Act would require location verification for exported chips, but the U.S. government doesn't currently have the infrastructure to enforce these controls - roughly 100,000 export-controlled GPUs were smuggled into China in 2024. Lucid is building verification infrastructure for AI training and deployment compliance without exposing the workload and data. The broader vision is to make all AI compute verifiable. Almost every compute governance proposal (export controls, international agreements, lab commitments) assumes the ability to verify what ran and what happened in any particular run. This infrastructure will be needed to enforce these proposals as they come into place.

• **Ulysses** - Secures AI datacenters against state-level attackers. Most datacenters used by labs sit around SL2-SL3 (using RAND's scale for weight security), and Ulysses is working on building SL4 datacenters from scratch, the level meant to hold against state-backed attack teams. This will be critical for when the labs inevitably train models with serious capacity for harm (ex. building WMDs).

• **Tinfoil** - Private inference that runs inside trusted execution environments (TEEs) on existing confidential-computing GPUs, keeping data provably private from the cloud provider, the host, and Tinfoil itself. This infrastructure is useful for independent model auditing, where an external party can run evals against a model without the lab seeing the test set or the auditor keeping the weights.

→ look into whether 3rd party data centers are popular right now / how many are founded

→ look into companies that keep users' data private - implying safety

Incentives & Institutions

Existing levers such as liability, insurance, and certification incentives can be used to drive safety standard adoption for companies selling and using AI tools, even while regulation is not yet in place. Companies

at this layer help define and drive adoption of AIS standards alongside policy. A widely adopted (or required) standard for AI agents, for example, gives companies an incentive to adopt best safety practices, which in turn creates a demand for services from companies at the other layers.

→ this and red-teaming are the easiest to prove most for, because they sort of have to be 3rd party

Highlights:

• **AIUC** - Creates a standard and certifies AI products for enterprise adoption - AIUC-1 is basically SOC-2 for agents, with insurance priced off the audits. The bet I like here is standard adoption: there's a market incentive to get certified even without regulation, and as regulation follows there's existing momentum and a widely adopted standard with clear implementation details. AIUC is already certifying agents for enterprise buyers, with ElevenLabs shipping agent insurance backed by AIUC-1.

• **Glacis** - Creates third-party verification for AI agents. They sell monitoring and steering software for deploying agents, but more critically they create an open standard for third-party verification which allows external parties to audit and insure AI use without sharing company data. There are existing mechanisms like liability for AI use, as well as explicit new policy mandates (Illinois SB 315, which takes effect in 2028) that create a market for third-party verification, and Glacis is building the infrastructure for it.

• **Exona Lab** - Helps price AI insurance - modeling AI risk so insurers can underwrite it. Pricing is the piece that's missing for AI insurance to exist at all, and once premiums move with how safely a company runs its AI, buying coverage starts paying for safety practices. They went through Lloyd's Lab in spring 2026, and their CEO co-authored a proposal for AI catastrophe bonds.

→ curious about the efficacy of this compared to AIUC's (unproven)

→ AIUC is backed by Lloyd's + recent push on mutual insurance - lack of insurance between labs (committing financially to safety)

Downstream Risks

As AIS approach (or have already reached) superhuman offensive bio and cyber capabilities, we have to secure all systems along the attack vectors that can currently be used to cause large-scale harm and destruction. Protein-screening, water supply monitoring, securing the electrical grid against cyber attacks, and creating tools for national security all arguably fall under the reducing-AI-risk thesis.

→ this and red-teaming are the easiest to prove most for, because they sort of have to be 3rd party

Highlights:

• **Vannevar Labs** - Builds AI-native open-source intelligence for US national security. Not a safety company by branding, but if you buy the MAM framing or the Total Research Transparency plan, deterrence and verification only work if states can see each other's AI programs, and someone has to build the collection layer that sees it. Vannevar sells exactly that for foreign-language open sources, profitable and at real scale.

• **Red Queen Bio** - Builds AI-powered bio-defense, aiming for countermeasures that exist before the pathogen does. AI is collapsing the cost of engineering pathogens faster than defenses improve, so they map what AI-enabled biotreatments would look like and pre-build the medical countermeasures against them. Spun out of HelixNano with a \$15M seed led by OpenAI.

Wildcard/Post-AI Transition

Even in a successful AI AGI transition without a catastrophic event, the post-AGI world opens up many new understudied problems. What does cognitive security look like when AIS are more persuasive than any human, and work is automated? How do we make sure all of humanity benefits from the AGI transformation? How do we keep human agency when delegating everything is the path of least resistance? This is the least clear and furthest-out area, but I think some of these problems are already salient and deserve attention now.

→ I'm offering orders for new chemical compounds etc. would fail lots

→ more people working on this will be proved in, so this becomes an important incentive to maintain

Highlight:

• **Workshop Labs** (joined Thinking Machines) - They were building user-aligned models - AIS aligned to the person using them rather than to the model provider. If all capable AI answers to a few labs, economic and political leverage concentrates with them, and everyone else becomes dependent on systems that don't represent them. Giving individuals models that are genuinely theirs keeps people economically relevant and agency distributed. They joined Thinking Machines in April 2026, and make the argument better than I can in their own essay.

Support Ecosystem

Beyond the startups themselves, there's an emerging ecosystem funding and incubating specifically for-profit AIS orgs.

Venture Capital

There's a handful of venture capital firms dedicated to funding AIS:

- **Lionheart Ventures** - first VC fund with a dedicated AI safety investing strategy, started in 2020. Transformative AI fund investing in AI security and assurance, pre-seed to Series A.
- **Juniper Ventures** - early-stage fund dedicated to AI assurance, pre-seed and seed with follow-ons (participated in Goodfire's Series B); has backed Goodfire, Gray Swan, Lucid, AIUC, and more.
- **Safe AI Fund (SAIF)** - \$100k checks in AIS startups.
- **Mythos Ventures** - early-stage fund for prosocial, safety-by-design AI startups; a broader portfolio with a few AIS bets (Goodfire, Softmax, AIUC).
- **Macroscopic Ventures** - catastrophic-risk focus, grants and equity - cooperative AI, misuse, AI welfare.

Incubators

There's a handful of incubators dedicated to funding AIS:

- **MATS** - Founding & Field-Building track - mentored by BlueDot, the SL5 Task Force, Halcyon Futures, and ARIA.
- **Halcyon Futures** - grants and mentorship for leaders at career inflection points building in AI safety, biosecurity, and AI security; 25+ projects launched. Also invests through Halcyon Ventures.
- **BlueDot Impact** - Incubator Week (up to \$100k in grants, no equity taken; the fifth cohort ran in late August 2026, with more likely to come) and **Rapid Grants**.
- **Seldon Lab** - an accelerator for AIS and security startups, with founder-friendly capital.
- **Constellation** - AIS incubator.
- **Coefficient Giving** - Navigators Incubator, a 12-week program under the Navigating Transformative AI fund; first cohort focused on securing transformative AI.
- **Surplus** - new incubator for software AIS startups, funding AI for epistemics & coordination, public-facing and community infra projects.
- **Institute for Progress (IFP)** - The Launch Sequence - matching proposals with founders.

The space is evolving rapidly, and the opportunities/markets develop faster as the public salience of the risks catches up and the Overton window shifts. In mid-2024 there were maybe a handful of companies you could call AIS startups, and now two years later there are dozens, with incubators and multiple VCs dedicated to supporting them. I think for prospective founders, there's a uniquely good opportunity to build in one of the emerging AIS markets. I plan to write up more thoughts on this soon.

Thanks for reading Yury's Substack!
Subscribe for free to receive new posts and support my work.